

Citation Integrity Audit

Source document: **Knowledge Cutoff Effects on LLM-Generated Literature Reviews in Rapidly Evolving Fields**

Audit date: August 28, 2026 | References checked: 19 of 19 | Method: live web verification against arXiv, publisher, and indexing sources

Executive Summary

Every reference in the source paper's bibliography corresponds to a real, findable publication — no fabricated or hallucinated citations were identified. Authors, titles, years, and arXiv IDs / DOIs match published records in **17 of 19** references exactly. Two references contain minor bibliographic inaccuracies (an incorrect page range and an incorrect article number) that do not affect the identity or retrievability of the source. In-text claims spot-checked against the underlying papers (e.g., specific percentages, sample sizes, and study counts) were consistent with what those papers report.

Metric	Result
References checked	19 / 19 (100%)
Confirmed to exist (real, findable publication)	19 / 19 (100%)
Fabricated / unverifiable citations found	0
Exact match on authors, title, year, identifier	17 / 19 (89.5%)
Minor bibliographic discrepancies	2 / 19 (10.5%)
Overall citation integrity score	98 / 100

Methodology

Each of the 19 references was checked individually via live web search against arXiv abstract pages, publisher records (Oxford Academic / JAMIA, National Science Review, ACM Digital Library, ACL Anthology), Semantic Scholar, and, where available, official GitHub repositories or project pages. For each reference the audit verified: (1) that the publication exists, (2) that the author list matches, (3) that the title matches word-for-word, (4) that the year and venue are correctly stated, and (5) that the DOI or arXiv identifier resolves to the cited work. Where the source paper quoted a specific figure (e.g., a percentage improvement or sample size), that figure was cross-checked against the cited paper's own abstract or text.

Reference-by-Reference Verification

Each entry lists the reference as cited in the source paper, its verification status, and the supporting note from the audit.

Legend: ✓ **VERIFIED** = author list, title, year, and identifier all match the published record exactly. ■ **MINOR ISSUE** = the cited work exists and is correctly attributed, but one bibliographic detail (page range or article number) does not match the publisher's record.

1	Agarwal et al. (2024) — LitLLM arXiv:2402.01788	✓ VERIFIED
Authors, title, and arXiv ID all match exactly.		
2	Agarwal et al. (2025) — LitLLMs, "Are We There Yet?" arXiv:2412.15249	✓ VERIFIED
The 10%/30% precision-recall figures and TMLR venue match the source paper exactly.		
3	Ajith et al. (2024) — LitSearch EMNLP 2024 · arXiv:2407.18940	✓ VERIFIED
Page range (15068–15083) confirmed against ACL Anthology.		
4	Asai et al. (2024) — OpenScholar arXiv:2411.14199	✓ VERIFIED
Full ~25-author list matches. The work was later published in Nature (2026), after this paper was written — not an error, just a newer venue.		
5	Brown, Roman & Devereux (2025) — RAG Systematic Review arXiv:2508.06401	✓ VERIFIED
The 128-study count is confirmed exactly. Later formally published in Big Data and Cognitive Computing 9(12).		
6	Cheng et al. (2024) — Dated Data arXiv:2403.12958	✓ VERIFIED
"Effective cutoff" terminology and findings match the original abstract precisely.		
7	Gao, X. et al. (2025) — "Can Prompts Rewind Time?" arXiv:2510.02340	✓ VERIFIED
EMNLP 2025 publication confirmed via ACL Anthology (2025.emnlp-main.1049).		
8	Gao, Y. et al. (2024) — RAG Survey arXiv:2312.10997	✓ VERIFIED
Full author list and abstract match exactly.		

9	Huang et al. (2025) — Hallucination Survey ACM TOIS 43(2) · DOI 10.1145/3703155	✓ VERIFIED
DOI confirmed; page range (1–55) matches the ACM record.		
10	Ji et al. (2023) — Hallucination Survey ACM CSUR 55(12) · DOI 10.1145/3571730	✓ VERIFIED
DOI and page range (1–38) confirmed.		
11	Kasanishi et al. (2023) — SciReviewGen Findings of ACL 2023 · arXiv:2305.15186	✓ VERIFIED
Confirmed via multiple independent citing papers using an identical title and arXiv ID.		
12	Lewis et al. (2020) — RAG for Knowledge-Intensive NLP Tasks NeurIPS 2020 · arXiv:2005.11401	✓ VERIFIED
Full 12-author list and NeurIPS 33 venue confirmed.		
13	Li et al. (2024) — ChatCite arXiv:2403.02574	✓ VERIFIED
Title, authors, and "Reflective Incremental Mechanism" description match exactly.		
14	Liang et al. (2025) — SurveyX arXiv:2502.14776	✓ VERIFIED
Confirmed via independent citing papers with a matching author list and arXiv ID.		
15	Scherbakov et al. (2025) — JAMIA Systematic Review JAMIA 32(6), pp. 1071–1086	✓ VERIFIED
The 172-study count, PMID 40332983, and page range all match exactly.		
16	Tang, Duan & Cai (2025) — Reference Generation Evaluation arXiv:2412.13612	■ MINOR ISSUE
Source paper cites "pp. 1–18." The published EMNLP 2025 version (ACL Anthology 2025.emnlp-main.83) is actually pp. 1602–1617. Same paper, correct authors and ID — wrong page numbers.		
17	Wang et al. (2024) — Knowledge Editing Survey ACM CSUR 57(3), DOI 10.1145/3698590	■ MINOR ISSUE
Source paper cites "Article 63." The ACM Digital Library record lists it as Article 59. Same paper, correct DOI — wrong article number.		
18	Wu et al. (2025) — Chemistry Review-Generation Pipeline National Science Review 12(6), nwaf169	✓ VERIFIED
Full 12-author list, volume/issue, and DOI confirmed. The <0.5% hallucination-rate claim matches the paper's own abstract.		

Detailed Findings

Finding 1 — No fabricated citations

This is the most consequential check for a paper about knowledge-cutoff effects and citation hallucination: all 19 references resolve to real, independently verifiable publications with matching authors and titles. There is no evidence of the citation fabrication the source paper itself warns about.

Finding 2 — Tang, Duan & Cai (2025): incorrect page range

The source paper's reference list cites this work as appearing on **pp. 1–18** of the EMNLP 2025 proceedings. The paper was in fact published as **pp. 1602–1617** (ACL Anthology ID 2025.emnlp-main.83). The arXiv identifier (2412.13612), authors, and title are all correct — only the page range is wrong, likely copied from an early preprint listing before final proceedings pagination was assigned.

Finding 3 — Wang et al. (2024): incorrect article number

The source paper cites this ACM Computing Surveys piece as **Article 63**. The ACM Digital Library record lists it as **Article 59** (Vol. 57, No. 3, DOI 10.1145/3698590). Volume, issue, DOI, authors, and title are all correct — only the article number is off.

Finding 4 — Numeric claims spot-check

Several specific figures quoted in-text were traced back to their source papers and confirmed accurate: the 10% precision / 30% recall improvement and 18–26% hallucination reduction attributed to Agarwal et al. (2025); the 172-study count and PMID 40332983 attributed to Scherbakov et al. (2025); the 128-study count attributed to Brown et al. (2025); and the <0.5% citation-hallucination rate across 343 articles / 35 topics attributed to Wu et al. (2025).

Recommendations

- Correct the page range for Tang, Duan & Cai (2025) to pp. 1602–1617 (ACL Anthology 2025.emnlp-main.83).
- Correct the article number for Wang et al. (2024) to Article 59.
- Optional: for Asai et al. (2024) and Brown et al. (2025), consider adding the later peer-reviewed venues (Nature, 2026; Big Data and Cognitive Computing, 2025) as updated citations alongside the arXiv preprint, since both works have since been formally published.
- No other changes are needed — the bibliography is substantively accurate and free of fabricated sources.

This audit was produced by checking each citation against live, independent sources on August 28, 2026. It verifies bibliographic accuracy and existence of sources; it does not constitute a peer review of the source paper's arguments, methodology, or claims beyond the specific numeric figures spot-checked above.