

Knowledge Cutoff Effects on LLM-Generated Literature Reviews in Rapidly Evolving Fields

A Review of Mechanisms, Current Mitigations, and Open Research Problems

Abstract

Large language models (LLMs) are increasingly used to draft, summarize, and even automate scientific literature reviews, yet every such model carries a fixed knowledge cutoff beyond which it has no direct exposure to newly published work. This paper examines how knowledge cutoff effects distort LLM-generated literature reviews specifically within rapidly evolving research fields, where publication velocity, terminological drift, and citation lag compress the effective shelf life of a model's parametric knowledge. Drawing on recent empirical work on effective versus reported cutoffs, temporal knowledge decay, citation hallucination, and retrieval-augmented generation (RAG), the paper synthesizes evidence that cutoff effects manifest not as a single sharp boundary but as a graded, resource-dependent degradation that produces recency bias, stale framing of “open problems,” omission of superseding work, and fabricated or outdated citations. Current mitigation strategies—retrieval augmentation, hybrid keyword–embedding search, plan-then-write generation pipelines, and post-hoc citation verification—reduce but do not eliminate these effects, particularly for the most recent months of a field's literature, where indexing lag and retrieval sparsity persist. The paper concludes by identifying research gaps around dynamic cutoff auditing, field-adaptive retrieval scheduling, and standardized benchmarks for time-sensitive review generation, and argues that treating knowledge cutoff as a first-class, quantifiable variable is necessary for the responsible use of LLMs in scholarly synthesis.

Keywords: large language models; knowledge cutoff; literature review generation; retrieval-augmented generation; citation hallucination; temporal generalization; scientific writing automation

1. Introduction

Literature reviews are among the most time-consuming yet foundational activities in scholarly research: they establish the boundary of known work, situate a contribution relative to prior art, and identify open problems. As the volume of scientific publication has grown, researchers have increasingly turned to large language models (LLMs) to accelerate literature discovery, summarization, and synthesis (Tang et al., 2025). At the same time, every LLM is trained on a corpus assembled up to a particular point in time — its knowledge cutoff — after which the model has no direct parametric awareness of newly published findings, terminology, or shifts in scientific consensus (Cheng et al., 2024).

This limitation is not merely inconvenient; it is potentially distorting when the object of the task is itself the state of a fast-moving literature. Fields such as large language model research, generative AI safety, and applied machine learning are characterized by publication cycles measured in weeks rather than years, rapid terminological churn, and frequent supersession of earlier findings. A literature review produced by an LLM whose training data ends several months or a year before the time of writing risks presenting an outdated “state of the art,” omitting work that has since superseded cited studies, or — in

the worst case — fabricating citations to plausible-sounding but non-existent papers when asked to describe recent developments it has not actually seen (Tang et al., 2025; Scherbakov et al., 2025).

This paper reviews the mechanisms, evidence, and mitigations relevant to knowledge cutoff effects in LLM-generated literature reviews, with particular attention to rapidly evolving fields. Section 2 surveys related work on knowledge cutoffs, temporal generalization, and automated review generation. Section 3 develops a technical account of how cutoff effects propagate into review quality. Section 4 surveys current mitigation approaches, principally retrieval augmentation and agentic pipelines. Section 5 discusses limitations of both the underlying phenomenon and the mitigation literature. Section 6 identifies research gaps and future directions, and Section 7 concludes.

2. Background and Related Work

2.1 The Knowledge Cutoff Construct

LLM developers typically report a single knowledge cutoff date alongside a released model, but empirical work shows this reported date is a simplification. Cheng et al. (2024) introduce the notion of an "effective cutoff," distinct from the designer-reported cutoff, and show through probing across dated versions of web-crawled resources that effective cutoffs vary by sub-resource and topic; they trace the discrepancy to temporal biases in Common Crawl snapshots, which retain substantial quantities of older content, and to deduplication procedures that inadvertently discard or conflate near-duplicate documents across time. The practical implication is that a model's knowledge of a given subfield does not end sharply at its official cutoff date but decays unevenly depending on how that subfield was represented, deduplicated, and refreshed in the training corpus.

Complementary evidence comes from work on temporally sensitive reasoning. Gao et al. (2025) show that prompting an LLM to simulate an earlier knowledge cutoff — asking it to reason as though it were situated at a past date — succeeds when the model is queried directly about information after that date, but fails to induce comparable forgetting of causally related downstream knowledge, indicating that cutoff boundaries are not cleanly enforced even when explicitly requested. This matters for review generation because authors sometimes prompt models to "only discuss work up to a given date" for reproducibility, yet such instructions do not reliably bound the model's actual knowledge use.

2.2 Hallucination and Factuality in LLM Outputs

Hallucination — the generation of plausible but unfounded or fabricated content — is a well-documented property of LLMs generally (Ji et al., 2023; Huang et al., 2025). Huang et al. (2025) provide a taxonomy distinguishing factuality hallucinations, where output contradicts verifiable real-world facts, from faithfulness hallucinations, where output is internally inconsistent with provided context or instructions; the survey further identifies limited or outdated parametric knowledge as one of the principal upstream causes of factuality hallucination. Citation fabrication is a specific and consequential instance of this phenomenon in academic contexts: a model asked to supply references for a claim may generate syntactically well-formed but non-existent bibliographic entries, particularly for topics at or beyond its training boundary.

2.3 Automated and LLM-Assisted Literature Review Generation

A growing body of systems targets literature review generation directly. LitLLM combines retrieval-augmented generation with LLM-based reranking to produce related-work sections from a query abstract (Agarwal et al., 2024), and a follow-up evaluation, Agarwal et al. (2025), decomposes the task into retrieval and plan-then-write generation, reporting that a hybrid of keyword-based and document-embedding-based search improves retrieval precision and recall over either method alone, and introduces a rolling evaluation protocol using recent arXiv papers specifically designed to limit test-set contamination for newly released models. ChatCite (Li et al., 2024) instead targets the summarization step, using a human-workflow-inspired agent with a reflective incremental mechanism to produce comparative literature summaries. OpenScholar (Asai et al., 2024) synthesizes scientific literature by grounding a retrieval-augmented LM in a large open corpus rather than relying solely on parametric knowledge, and LitSearch (Ajith et al., 2024) provides a dedicated benchmark for scientific literature retrieval, an upstream capability on which review quality depends. Wu et al. (2025) report an end-to-end pipeline applied to a chemistry subfield (propane dehydrogenation catalysis) that, through multilayered quality control including expert verification, reduced citation hallucination to under 0.5% with 95% confidence, illustrating that careful engineering can substantially curb — though not fully eliminate — the problem in a bounded domain.

Two systematic evaluations are particularly relevant to the present paper. Tang et al. (2025) introduce a multidimensional framework assessing hallucination rates in generated references alongside semantic coverage and factual consistency of generated summaries and reviews against human-written counterparts; they find that even the most advanced models evaluated continued to generate hallucinated references, and that performance varied by discipline. Scherbakov et al. (2025), a large language model-assisted systematic review of 172 eligible studies on LLM use in review automation, similarly report that literature search and citation generation exhibited comparatively high hallucination rates relative to other review stages such as screening and data extraction, and note that this was especially pronounced for earlier GPT-based model generations. Neither study isolates the knowledge cutoff variable per se, but both provide corroborating evidence that the citation-and-currency dimension of review generation is disproportionately fragile compared with other subtasks such as classification or extraction from provided text.

3. Mechanisms of Knowledge Cutoff Effects on Literature Reviews

This section develops a more granular account of how a fixed knowledge cutoff propagates into specific failure modes of a generated literature review, organized around four mechanisms: uneven temporal decay, recency-anchored framing, citation and reference fabrication, and terminological drift.

3.1 Uneven Temporal Decay Rather Than a Sharp Boundary

A naive model of knowledge cutoff treats it as a wall: everything before the date is known, everything after is unknown. Empirical work contradicts this. Cheng et al. (2024) find effective cutoffs differing substantially from reported cutoffs at the resource level, driven by uneven representation of dated content in the pretraining corpus. Related evidence outside the review-generation literature indicates

that even within a domain, model recall of "current" knowledge just before a cutoff already shows degradation, and that recall of superseded, historical information is often weaker than recall of the most recent consensus — suggesting an asymmetric forgetting pattern rather than uniform retention up to a hard boundary. For literature review generation, the practical consequence is that a model's coverage of a rapidly evolving subfield is not merely truncated at a known date; it can be sparse or biased for the months immediately preceding the nominal cutoff, precisely the window in which the newest and most citation-relevant work would have appeared.

3.2 Recency-Anchored Framing of "Open Problems"

Literature reviews conventionally close with a discussion of open problems and future directions. When generated by an LLM, this section is especially vulnerable to cutoff effects because it requires the model to assess what remains unsolved as of the present moment — a moment the model cannot directly observe. In rapidly evolving fields, problems flagged as "open" at a model's training cutoff are frequently addressed, reframed, or rendered moot within months. A review that reproduces a cutoff-era framing of open problems therefore risks presenting a picture of the field's frontier that is materially stale by the time a human reader consults it, even if every individual factual claim about pre-cutoff work is accurate.

3.3 Citation and Reference Fabrication Under Extrapolation Pressure

When a user's prompt implies coverage of the current literature but the model's knowledge does not extend that far, the model faces what might be termed extrapolation pressure: an instruction to describe content it has not seen. Under this pressure, LLMs have been repeatedly documented to fabricate citations — generating references with plausible author names, journal titles, and years that do not correspond to real publications (Tang et al., 2025; Scherbakov et al., 2025; Huang et al., 2025). This is distinct from, but compounded by, the general hallucination tendency of LLMs; the knowledge cutoff specifically determines the point past which any claim of citing a "recent" paper is necessarily either retrieved from an external source or fabricated, since it cannot come from unaided parametric recall.

3.4 Terminological and Conceptual Drift

Fast-moving fields frequently rename concepts, consolidate previously distinct subareas, or deprecate earlier terminology in favor of new consensus vocabulary. A model trained before such a shift may use outdated terms, fail to recognize that a newer term denotes largely the same construct as one it knows, or miscategorize a paper's contribution relative to a taxonomy that has since been revised. This drift is difficult to detect through simple factual-accuracy checks, because individual sentences may be accurate relative to the terminology current at the model's cutoff while collectively producing a review that reads as anachronistic to a domain expert.

4. Current Approaches and Developments

4.1 Retrieval-Augmented Generation

The dominant mitigation strategy is retrieval-augmented generation (RAG), in which a generative model is grounded in documents retrieved from an external, updatable corpus at inference time rather than

relying solely on parametric memory (Lewis et al., 2020). Subsequent surveys describe RAG as directly addressing two coupled weaknesses of LLMs: hallucination and outdated internal knowledge (Gao et al., 2024), and a recent systematic review of 128 highly cited RAG studies published between 2020 and mid-2025 confirms that grounding generation in non-parametric, retrievable evidence is now the standard architectural response to knowledge staleness in LLM applications (Brown et al., 2025). Applied to literature review generation, RAG in principle decouples review currency from model training cutoff: as long as the retrieval index is refreshed, the system can surface papers published after the underlying LLM's cutoff, with the model's role limited to synthesis and summarization of retrieved content rather than recall from memory.

In practice, this decoupling is only partial. Retrieval quality depends on the completeness and recency of the underlying index, which is itself subject to indexing lag on preprint servers and publisher platforms. Agarwal et al. (2025) report that combining keyword-based and document-embedding-based search substantially improves precision and recall relative to either method alone — by roughly 10% and 30%, respectively, in their reported evaluation — but also find that even their best retrieval configuration does not achieve complete recall of relevant recent work, and note that plan-based generation (first outlining a review structure, then executing retrieval and writing against that plan) further reduces but does not eliminate hallucinated content relative to single-pass generation.

4.2 Agentic and Multi-Stage Pipelines

Beyond single-pass RAG, several systems decompose review writing into discrete agentic stages mirroring human workflows. ChatCite's Key Element Extractor and Reflective Incremental Generator explicitly separate information extraction from comparative synthesis, aiming to reduce the compounding of retrieval and generation errors (Li et al., 2024). OpenScholar similarly grounds synthesis in retrieved passages from a large scientific corpus rather than treating generation as an unconstrained parametric task (Asai et al., 2024). These architectures share a common logic: by making the retrieval step explicit, verifiable, and separable from generation, they create a natural point at which recency and provenance can be audited before synthesis occurs.

4.3 Domain-Specific Quality Control Pipelines

In bounded domains, multilayered quality-control pipelines have demonstrated that hallucination — including citation fabrication tied to currency gaps — can be driven very low. Wu et al. (2025) report a domain-specific system for chemistry review generation combining automated verification steps with expert review, achieving citation-integrity hallucination rates below 0.5% at 95% confidence across reviews spanning 343 source articles and 35 topics. This suggests that the residual cutoff-and-hallucination problem is tractable when a system can afford extensive verification and a relatively bounded, well-indexed corpus, but raises the open question of whether comparable rigor is achievable at the speed and breadth demanded by genuinely rapidly evolving fields such as generative AI research itself, where the literature under review is simultaneously the domain producing the LLMs used to review it.

4.4 Benchmarking Recency-Sensitive Generation

Recognizing that static benchmarks are quickly contaminated by newer training runs, Agarwal et al. (2025) propose a rolling evaluation protocol that constructs test sets from arXiv papers on a recurring basis, explicitly designed to remain valid for newly released LLMs by sampling papers unlikely to have been seen during training. This class of evaluation design is a direct methodological response to the knowledge-cutoff problem: rather than assuming a fixed, contamination-free test set, it treats recency as a moving target that must be continuously re-established relative to each model's training boundary.

5. Research Challenges and Limitations

5.1 Retrieval Does Not Fully Substitute for Currency

Even fully retrieval-augmented pipelines inherit the recency limits of their retrieval index. Preprint servers, journal publication pipelines, and citation databases each have their own latency between a work's dissemination and its appearance in a searchable, well-indexed form. A RAG system can therefore still under-represent the most recent weeks or months of a fast-moving field's literature even when its underlying LLM's parametric cutoff is not the binding constraint — a distinction that is easy to elide when reporting a system as "solving" the knowledge-cutoff problem.

5.2 Reported Cutoffs Are Unreliable Proxies for Actual Knowledge

Because effective cutoffs vary by sub-resource and topic (Cheng et al., 2024), a single reported cutoff date cannot be used as a reliable guarantee of a model's knowledge boundary for a specific subfield. This complicates both user expectations and evaluation design: a benchmark constructed around a model's official cutoff date may inadvertently test topics for which the model's effective cutoff is earlier or later than reported, confounding measurement of cutoff-related failure rates.

5.3 Persistent Hallucination Under Extrapolation Demand

Across evaluations, hallucination in reference generation persists even in advanced models and even when retrieval is available, with disciplinary variation in severity (Tang et al., 2025). This indicates that citation fabrication in the currency-critical portion of a review is not simply an artifact of missing retrieval infrastructure; it also reflects a generation-time tendency to produce fluent, well-formed but unverified content when a task implicitly demands coverage beyond what has been reliably retrieved or verified.

5.4 Prompted Cutoff Simulation Is an Unreliable Control

A tempting but incomplete safeguard is to explicitly instruct a model to restrict itself to information available before a stated date, whether to enforce a review's stated scope or to construct a contamination-free evaluation. Gao et al. (2025) show that such prompted cutoffs work reasonably well when a model is asked directly about post-cutoff facts but fail to suppress use of causally related downstream knowledge, meaning a model may still leak influence from later developments into a review nominally scoped to an earlier date. This limits the reliability of prompt-only strategies for both authorial scoping and rigorous evaluation design.

5.5 Evaluation Methodology Gaps

Much of the evaluation literature on LLM-generated reviews measures hallucination rate, semantic coverage, or factual consistency against a fixed human-written comparator (Tang et al., 2025), but comparatively little work isolates knowledge cutoff as an independent variable — for instance, by systematically varying the gap between a model's training cutoff and the target review's intended "as-of" date, or by decomposing hallucination rates according to whether the cited work predates or postdates the model's effective cutoff. Without such decomposition, it remains difficult to attribute observed hallucination or omission specifically to cutoff-driven staleness versus other causes of LLM unreliability.

6. Research Gaps and Future Directions

6.1 Field-Adaptive Retrieval Scheduling

Rapidly evolving fields would benefit from retrieval scheduling policies that adapt update frequency and search breadth to a field's empirically measured publication velocity, rather than treating all domains as equally slow-moving. Building on the hybrid keyword–embedding retrieval strategies shown to improve recall (Agarwal et al., 2025), future systems could incorporate explicit recency-weighting or field-velocity signals that widen retrieval windows and increase index refresh frequency specifically for subfields identified as fast-moving.

6.2 Dynamic, Sub-Resource-Level Cutoff Auditing

Given that effective cutoffs vary by topic and resource (Cheng et al., 2024), a promising direction is the development of lightweight, user-facing tools that estimate a model's effective — rather than reported — cutoff for a specific subfield before a review is generated, allowing authors and readers to calibrate trust in the parametric portion of a review's coverage on a per-topic basis rather than relying on a single global cutoff date.

6.3 Cutoff-Decomposed Evaluation Benchmarks

Extending rolling, contamination-resistant evaluation protocols such as that of Agarwal et al. (2025), future benchmarks could explicitly stratify test items by their publication date relative to each evaluated model's cutoff, enabling direct measurement of how hallucination rate, omission rate, and framing staleness scale with the pre-/post-cutoff time gap, rather than reporting a single aggregate hallucination figure that conflates cutoff-driven and non-cutoff-driven errors.

6.4 Provenance-Transparent Review Generation

Systems could be designed to explicitly tag each claim or citation in a generated review with its source — retrieved-and-verified, retrieved-but-unverified, or parametric recall — surfacing to the reader precisely which portions of a review are grounded in externally verifiable, dated evidence versus the model's internal, cutoff-bounded knowledge. This would extend the provenance orientation already present in citation-aware retrieval systems such as OpenScholar (Asai et al., 2024) toward an explicit, per-claim currency label.

6.5 Cross-Domain Generalization of Quality-Control Pipelines

The multilayered quality-control approach that achieved sub-0.5% hallucination in a bounded chemistry domain (Wu et al., 2025) warrants systematic testing in domains with substantially higher publication velocity and more contested, less consolidated terminology — most notably AI and machine learning research itself — to determine whether comparable rigor can be sustained without prohibitive verification cost as publication rate increases.

7. Conclusion

Knowledge cutoff is frequently discussed as a simple date beyond which an LLM "knows nothing," but the evidence reviewed here indicates a more complex reality: effective cutoffs vary by topic and resource (Cheng et al., 2024), prompted attempts to simulate an earlier cutoff only partially suppress downstream knowledge leakage (Gao et al., 2025), and the citation-and-currency dimension of literature review generation remains disproportionately prone to hallucination relative to other subtasks even in state-of-the-art systems (Tang et al., 2025; Scherbakov et al., 2025). Retrieval-augmented and agentic pipelines meaningfully mitigate — but do not eliminate — these effects, and their residual failure modes are concentrated precisely in the most recent slice of a field's literature, which is also the slice most consequential for accurately characterizing a rapidly evolving field's current state. Treating knowledge cutoff as a quantifiable, sub-resource-varying, and independently measurable property of LLM-generated scholarly text — rather than a single reported date to be mentioned in passing — is a necessary step toward literature reviews that researchers in fast-moving fields can responsibly rely upon.

References

- Agarwal, S., Laradji, I. H., Charlin, L., & Pal, C. (2024). LitLLM: A toolkit for scientific literature review. arXiv preprint arXiv:2402.01788. <https://doi.org/10.48550/arXiv.2402.01788>
- Agarwal, S., Sahu, G., Puri, A., Laradji, I. H., Dvijotham, K. D. J., Stanley, J., Charlin, L., & Pal, C. (2025). LitLLMs, LLMs for literature review: Are we there yet? Transactions on Machine Learning Research. arXiv preprint arXiv:2412.15249. <https://doi.org/10.48550/arXiv.2412.15249>
- Ajith, A., Xia, M., Chevalier, A., Goyal, T., Chen, D., & Gao, T. (2024). LitSearch: A retrieval benchmark for scientific literature search. In Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing (pp. 15068–15083). Association for Computational Linguistics. arXiv:2407.18940. <https://doi.org/10.48550/arXiv.2407.18940>
- Asai, A., He, J., Shao, R., Shi, W., Singh, A., Chang, J. C., Lo, K., Soldaini, L., Feldman, S., D'Arcy, M., et al. (2024). OpenScholar: Synthesizing scientific literature with retrieval-augmented LMs. arXiv preprint arXiv:2411.14199. <https://doi.org/10.48550/arXiv.2411.14199>
- Brown, A., Roman, M., & Devereux, B. (2025). A systematic literature review of retrieval-augmented generation: Techniques, metrics, and challenges. arXiv preprint arXiv:2508.06401. <https://doi.org/10.48550/arXiv.2508.06401>
- Cheng, J., Marone, M., Weller, O., Lawrie, D., Khashabi, D., & Van Durme, B. (2024). Dated data: Tracing knowledge cutoffs in large language models. arXiv preprint arXiv:2403.12958. <https://doi.org/10.48550/arXiv.2403.12958>

- Gao, X., Zhang, R., Du, D., Mahindre, S., Somayajula, S. A., & Xie, P. (2025). Can prompts rewind time for LLMs? Evaluating the effectiveness of prompted knowledge cutoffs. In Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing. arXiv:2510.02340. <https://doi.org/10.48550/arXiv.2510.02340>
- Gao, Y., Xiong, Y., Gao, X., Jia, K., Pan, J., Bi, Y., Dai, Y., Sun, J., Wang, M., & Wang, H. (2024). Retrieval-augmented generation for large language models: A survey. arXiv preprint arXiv:2312.10997. <https://doi.org/10.48550/arXiv.2312.10997>
- Huang, L., Yu, W., Ma, W., Zhong, W., Feng, Z., Wang, H., Chen, Q., Peng, W., Feng, X., Qin, B., & Liu, T. (2025). A survey on hallucination in large language models: Principles, taxonomy, challenges, and open questions. ACM Transactions on Information Systems, 43(2), 1–55. <https://doi.org/10.1145/3703155>
- Ji, Z., Lee, N., Frieske, R., Yu, T., Su, D., Xu, Y., Ishii, E., Bang, Y. J., Madotto, A., & Fung, P. (2023). Survey of hallucination in natural language generation. ACM Computing Surveys, 55(12), 1–38. <https://doi.org/10.1145/3571730>
- Kasanishi, T., Isonuma, M., Mori, J., & Sakata, I. (2023). SciReviewGen: A large-scale dataset for automatic literature review generation. In Findings of the Association for Computational Linguistics: ACL 2023. arXiv:2305.15186. <https://doi.org/10.48550/arXiv.2305.15186>
- Lewis, P., Perez, E., Piktus, A., Petroni, F., Karpukhin, V., Goyal, N., Küttler, H., Lewis, M., Yih, W., Rocktäschel, T., Riedel, S., & Kiela, D. (2020). Retrieval-augmented generation for knowledge-intensive NLP tasks. In Advances in Neural Information Processing Systems 33 (NeurIPS 2020). arXiv:2005.11401. <https://doi.org/10.48550/arXiv.2005.11401>
- Li, Y., Chen, L., Liu, A., Yu, K., & Wen, L. (2024). ChatCite: LLM agent with human workflow guidance for comparative literature summary. arXiv preprint arXiv:2403.02574. <https://doi.org/10.48550/arXiv.2403.02574>
- Liang, X., Yang, J., Wang, Y., Tang, C., Zheng, Z., Song, S., Lin, Z., Yang, Y., Niu, S., Wang, H., et al. (2025). SurveyX: Academic survey automation via large language models. arXiv preprint arXiv:2502.14776. <https://doi.org/10.48550/arXiv.2502.14776>
- Scherbakov, D., Hubig, N., Jansari, V., Bakumenko, A., & Lenert, L. A. (2025). The emergence of large language models as tools in literature reviews: A large language model-assisted systematic review. Journal of the American Medical Informatics Association, 32(6), 1071–1086. <https://doi.org/10.1093/jamia/ocaf063>. PMID: 40332983.
- Tang, X., Duan, X., & Cai, Z. G. (2025). Large language models for automated literature review: An evaluation of reference generation, abstract writing, and review composition. In Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing (pp. 1–18). Association for Computational Linguistics. arXiv:2412.13612. <https://doi.org/10.48550/arXiv.2412.13612>
- Wang, S., Zhu, Y., Liu, H., Zheng, Z., Chen, C., & Li, J. (2024). Knowledge editing for large language models: A survey. ACM Computing Surveys, 57(3), Article 63. <https://doi.org/10.1145/3698590>
- Wu, S., Ma, X., Luo, D., Li, L., Shi, X., Chang, X., Lin, X., Luo, R., Pei, C., Du, C., Zhao, Z.-J., & Gong, J. (2025). Automated literature research and review-generation method based on large language models. National Science Review, 12(6), nwaf169. <https://doi.org/10.1093/nsr/nwaf169>

Yao, Y., Wang, P., Tian, B., Cheng, S., Li, Z., Deng, S., Chen, H., & Zhang, N. (2023). Editing large language models: Problems, methods, and opportunities. In Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing (pp. 10222–10240). Association for Computational Linguistics. <https://doi.org/10.18653/v1/2023.emnlp-main.632>